Machine Translation

Tower of Babel

Google Translate

Text | Documents

DETECT LANGUAGE | ENGLISH | SPANISH | FRENCH

ENGLISH | SPANISH | ARABIC

Search languages

Detect language	Danish	Hmong	Lithuanian	Romanian	Telugu
Afrikaans	Dutch	Hungarian	Luxembourgish	Russian	Thai
Albanian	English	Icelandic	Macedonian	Samoan	Turkish
Amharic	Esperanto	Igbo	Malagasy	Scots Gaelic	Turkmen
Arabic	Estonian	Indonesian	Malay	Serbian	Ukrainian
Armenian	Filipino	Irish	Malayalam	Sesotho	Urdu
Azerbaijani	Finnish	Italian	Maltese	Shona	Uyghur
Basque	French	Japanese	Maori	Sindhi	Uzbek
Belarusian	Frisian	Javanese	Marathi	Sinhala	Vietnamese
Bengali	Galician	Kannada	Mongolian	Slovak	Welsh
Bosnian	Georgian	Kazakh	Myanmar (Burmese)	Slovenian	Xhosa
Bulgarian	German	Khmer	Nepali	Somali	Yiddish
Catalan	Greek	Kinyarwanda	Norwegian	Spanish	Yoruba
Cebuano	Gujarati	Korean	Odia (Oriya)	Sundanese	Zulu

Dictionaries

- English: leg, foot, paw
- French: jambe, pied, patte, etape

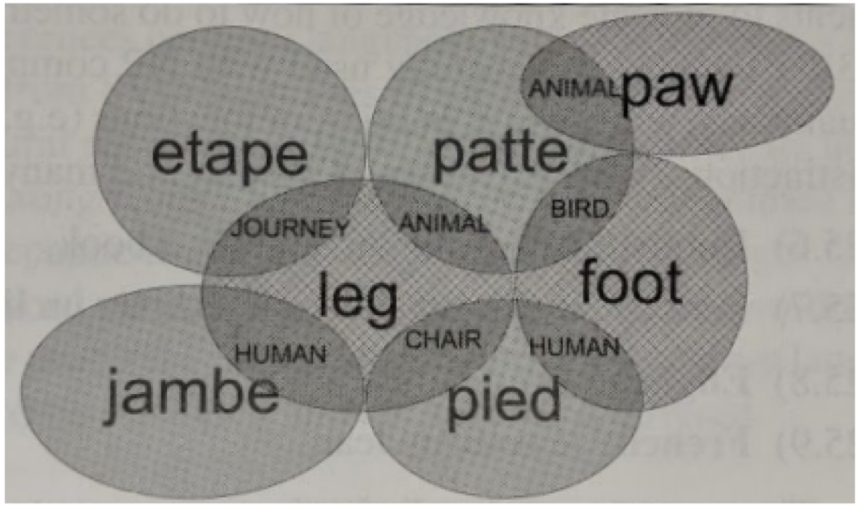

Challenges

- Ambiguities
  - Words
  - Morphology
  - Syntax
  - Semantics
  - Pragmatics
- Gaps in data
  - Availability of corpus
  - Commonsense knowledge
- Understanding of context, connotation, social norms, etc

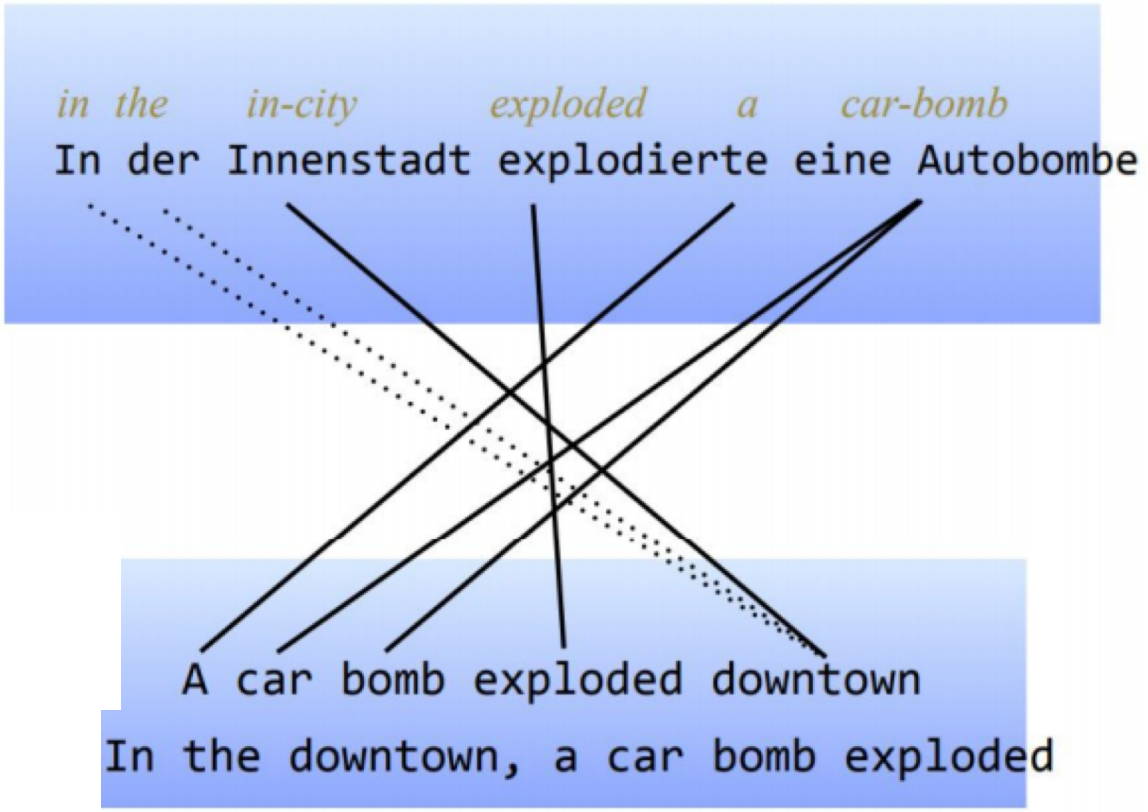

Research Problems

- How can we formalize the process of learning to translate from examples?
- How can we formalize the process of finding translations for new inputs?
- If our model produces many outputs, how do we find the best one?
- If we have a gold standard translation, how can we tell if our output is good or bad?

Two Views Of MT

MT as Code Breaking

One naturally wonders if the problem of translation could conceivably be treated as a problem in cryptography. When I look at an article in Russian, I say: '***This is really written in English, but it has been coded in some strange symbols. I will now proceed to decode.***'

Warren Weaver to Norbert Wiener, March, 1947

The Noisy-Channel Model

Claude Shannon. "A Mathematical Theory of Communication" 1948.

The Noisy-Channel Model

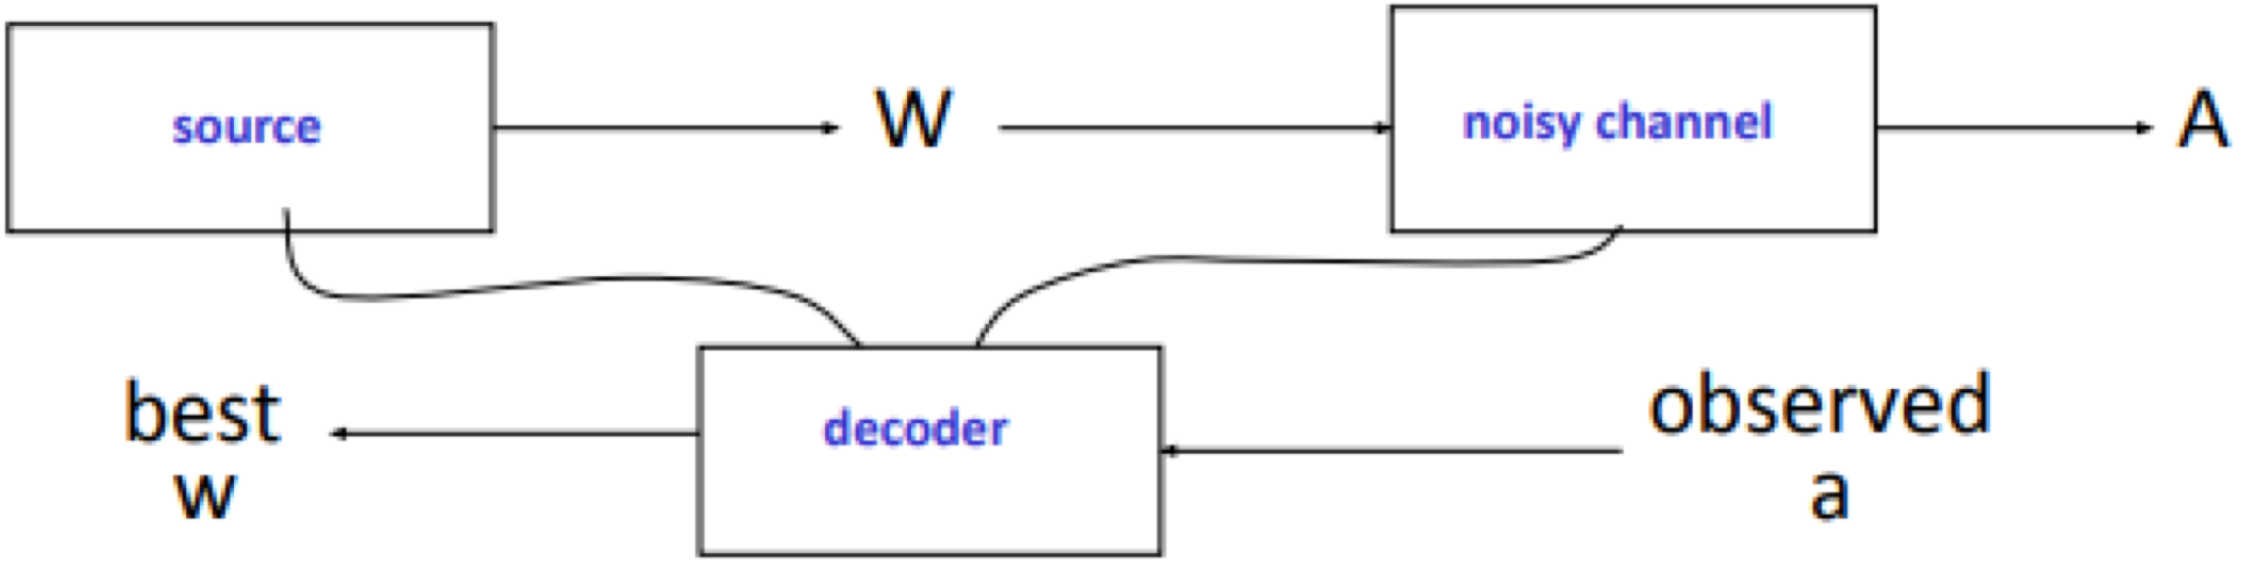

The Noisy-Channel Model

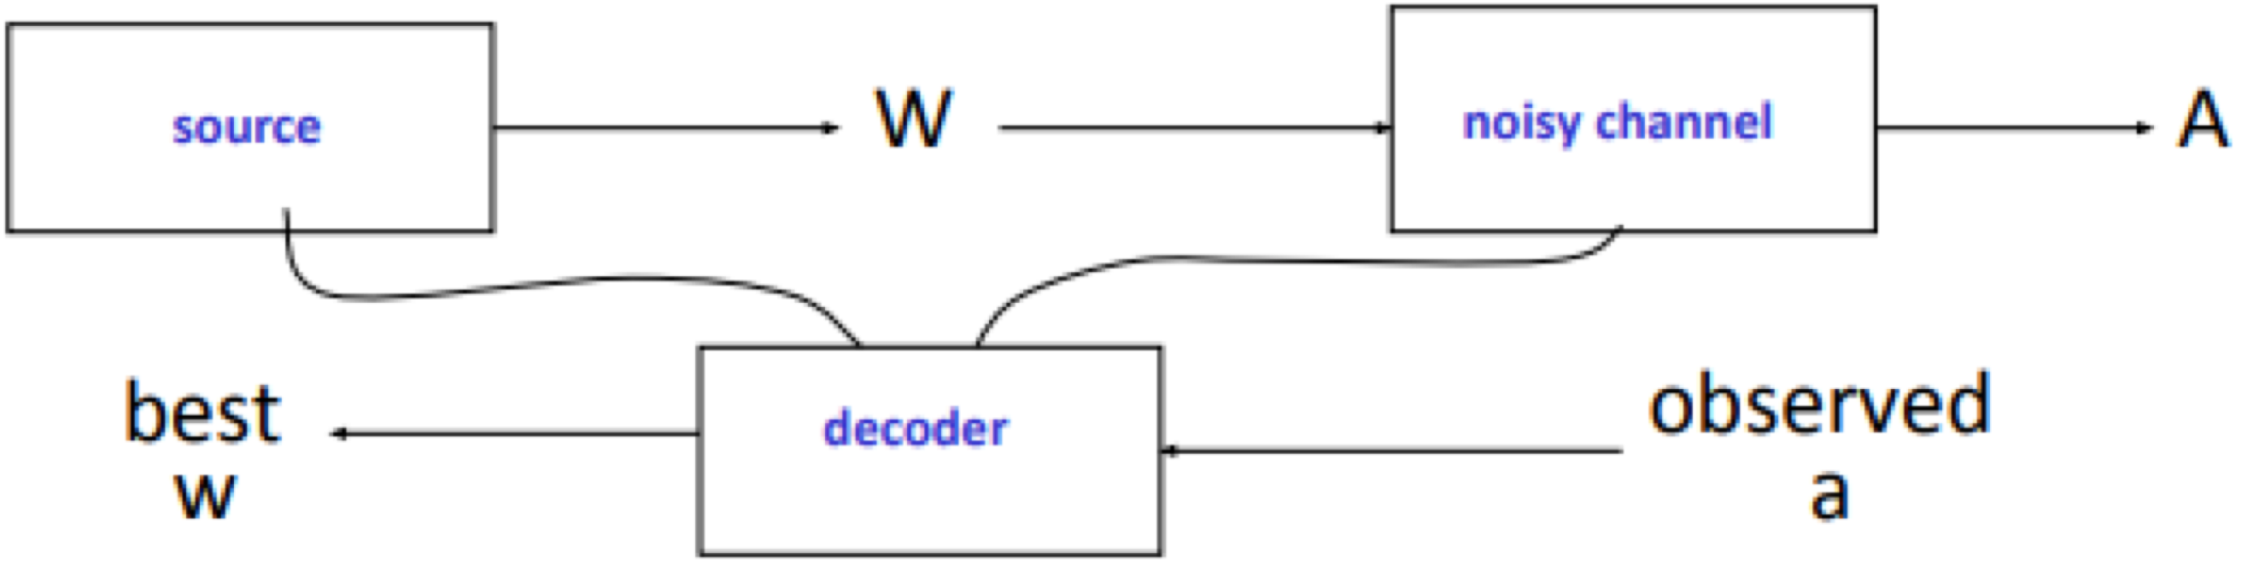

We want to predict a sentence given acoustics:

$$w^* = \arg\max_{w} P(w|a)$$

The Noisy-Channel Model

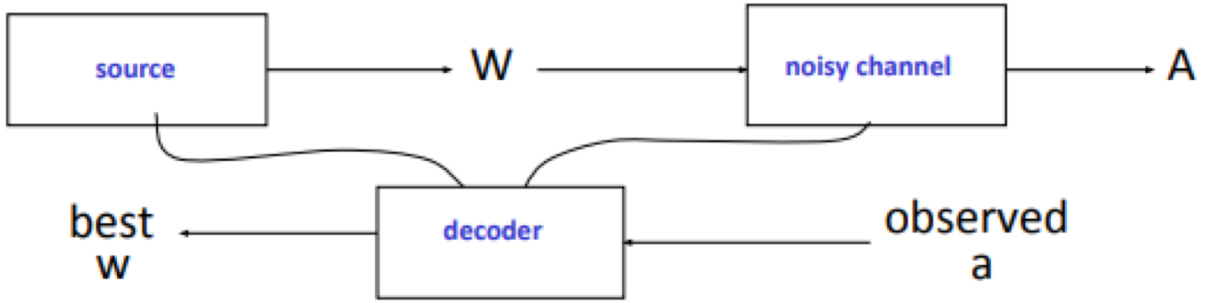

$$w^* = \arg\max_{w} P(w|a)$$

$$= \arg\max_{w} \boldsymbol{P(a|w)}\ \boldsymbol{P(w)} \ / \ P(a)$$

$$= \arg\max_{w} \boldsymbol{P(a|w)}\ \boldsymbol{P(w)}$$

Channel model

Source model

The Noisy-Channel Model

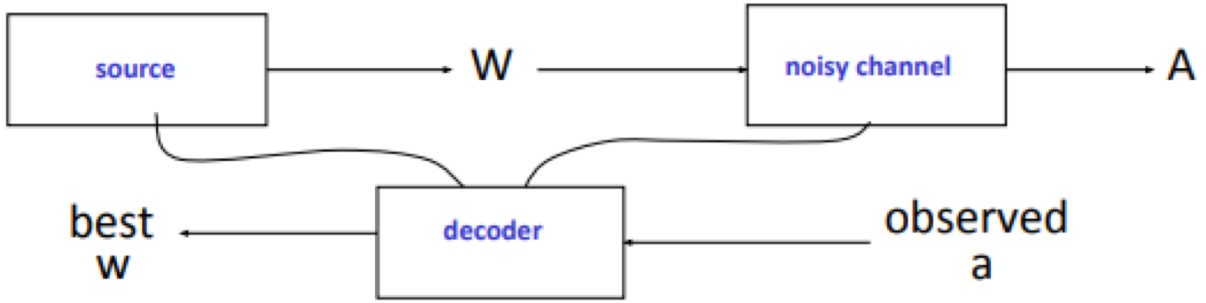

$$w^* = \arg\max_{w} P(w|a)$$

$$= \arg\max_{w} \boldsymbol{P(a|w)}\ \boldsymbol{P(w)} \ / \ P(a)$$

$$= \arg\max_{w} \boldsymbol{P(a|w)}\ \boldsymbol{P(w)}$$

Likelihood
Acoustic model (HMMs)
Translation model

Prior
Language model: Distributions over sequence of words

The Noisy-Channel Model

$$\hat{e} = \arg\max_{e} p_{\varphi}(e) \times p_{\theta}(f \mid e)$$

Language model

Translation model

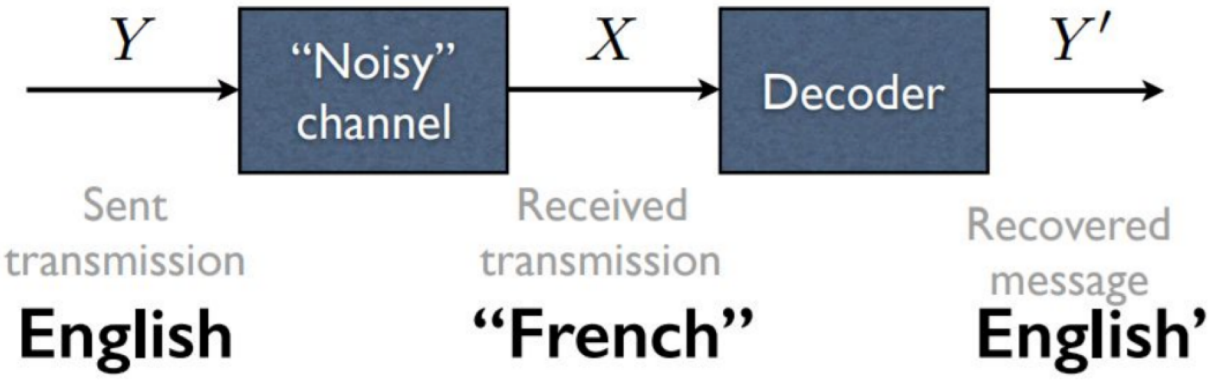

MT as Direct Modeling

$$\hat{e} = \arg\max_{e} p_\lambda(e \mid f)$$

target

source

- One model does everything
- Trained to reproduce a corpus of translations

Two Views of MT

- Code breaking (aka the noisy channel, Bayes rule)
  - I know the target language
  - I have example translations texts (example enciphered data)
- Direct modeling (aka pattern matching)
  - I have really good learning algorithms and a bunch of example inputs (source language sentences) and outputs (target language translations)

Which is Better?

- **Noisy channel** - $p_{\varphi}(e) \times p_{\theta}(f \mid e)$
  - Easy to use monolingual target language data
  - Search happens under a product of two models (individual models can be simple, product can be powerful)
- **Direct Model** - $p_{\lambda}(e \mid f)$
  - Directly model the process you care about
  - Model must be very powerful

Where are we in 2020?

- Direct modeling is where most of the action is
  - Neural networks are very good at generalizing and conceptually very simple
  - Inference in “product of two models” is hard
- Noisy channel ideas are incredibly important and still play a big role in how we think about translation

Two Views of MT

- Noisy channel $\hat{e} = \arg\max_{e} p_{\varphi}(e) \times p_{\theta}(f \mid e)$
- Direct $\hat{e} = \arg\max_{e} p_{\lambda}(e \mid f)$

Noisy Channel: Phrase-Based MT

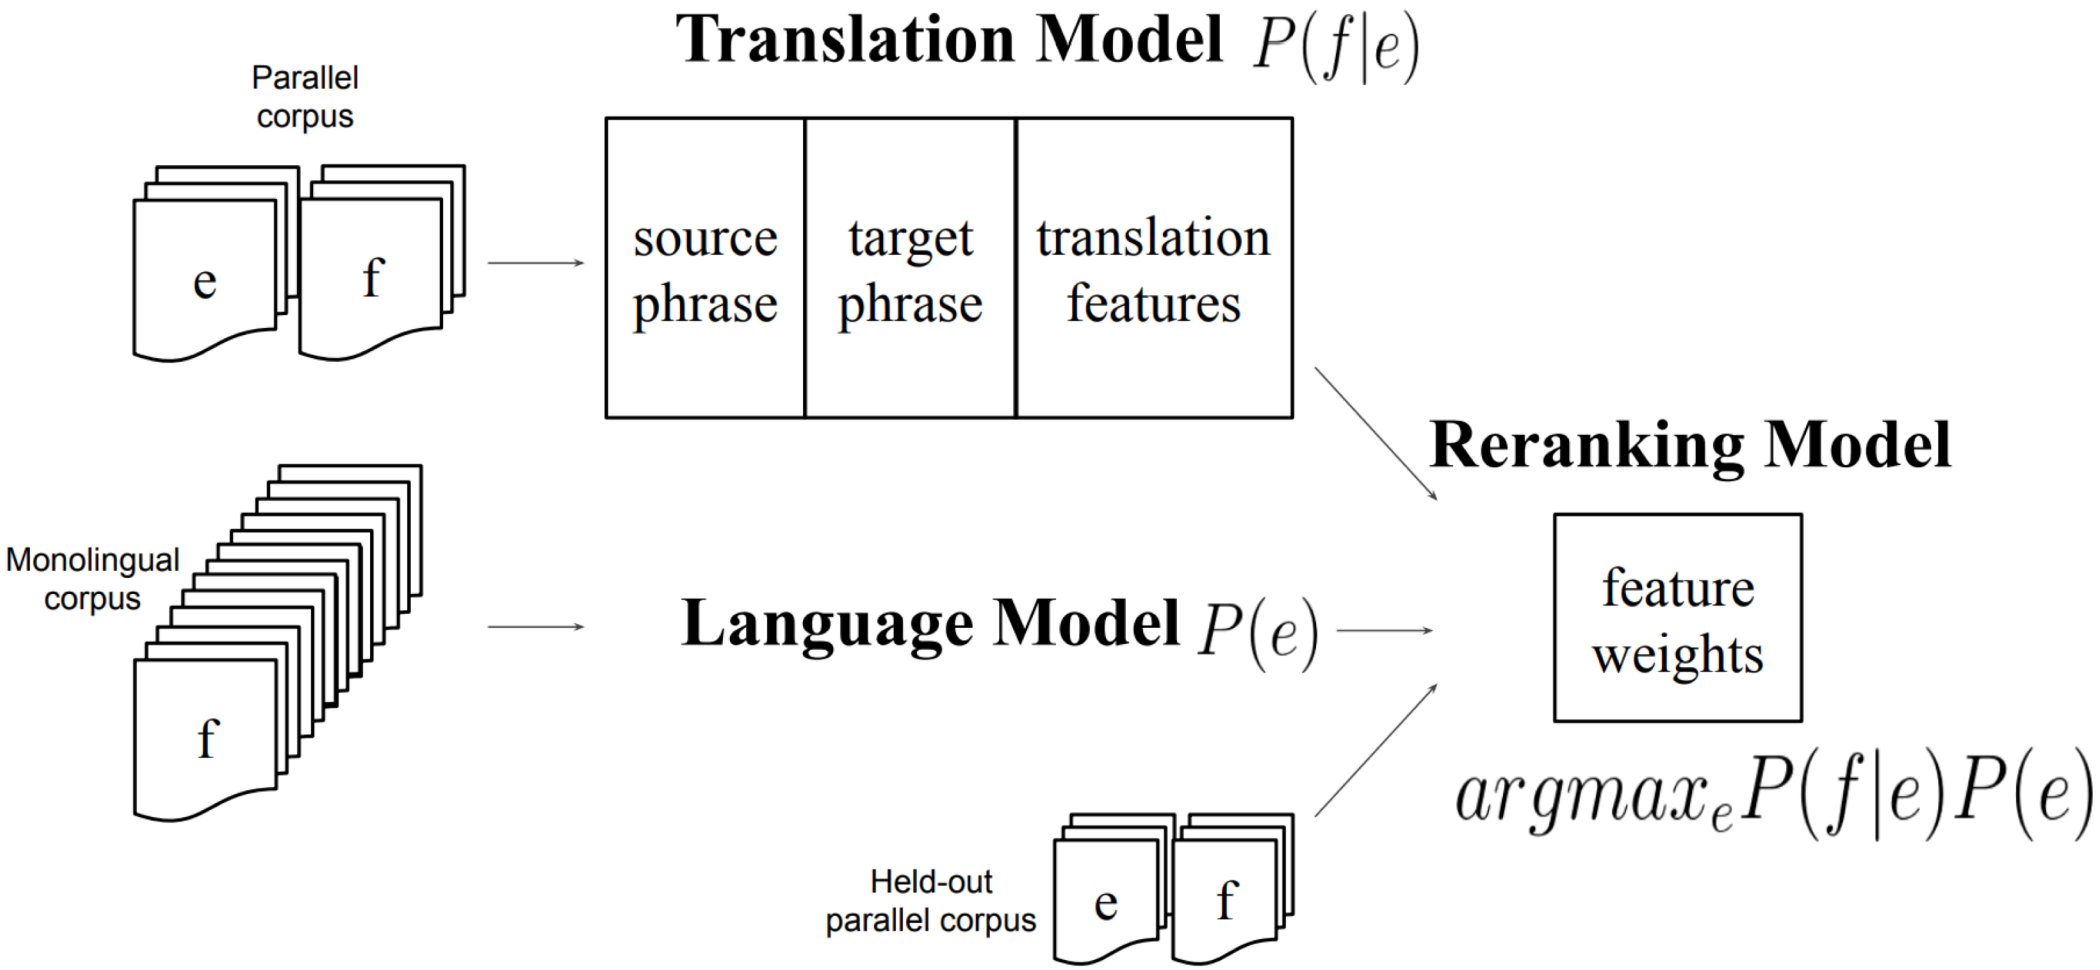

Neural MT: Conditional Language Modeling

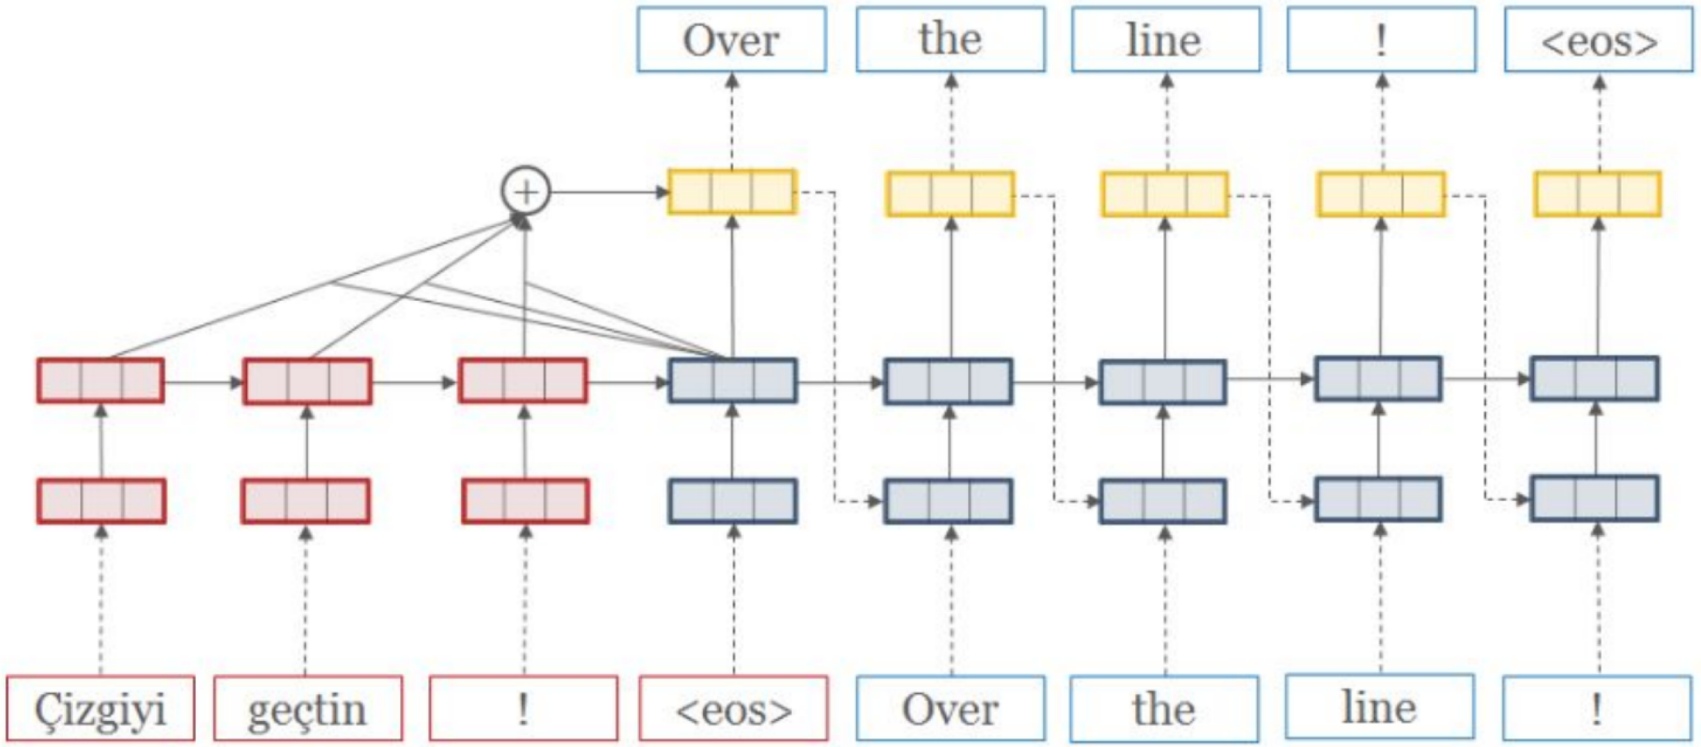

http://opennmt.net/

A Common Problem

- Noisy channel $\hat{e} = \arg\max_{e} p_{\varphi}(e) \times \boldsymbol{p_{\theta}(f \mid e)}$
- Direct $\hat{e} = \arg\max_{e} \boldsymbol{p_{\lambda}(e \mid f)}$

- Both models must assign probabilities to how a sentence in one language translates into a sentence in another language

Learning From Data

Parallel Corpora

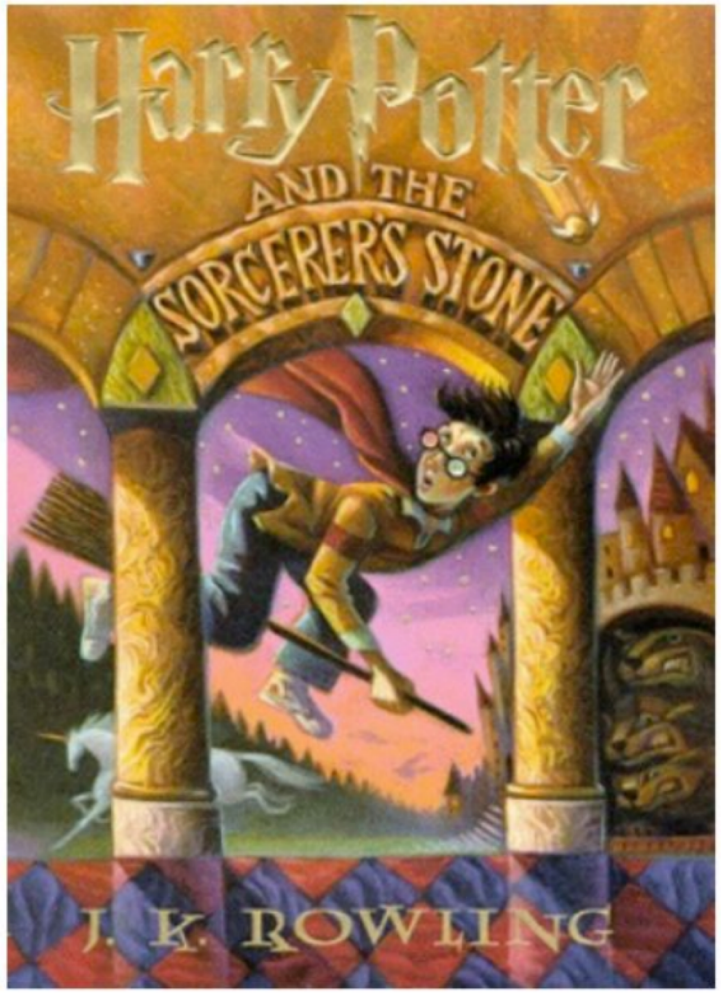

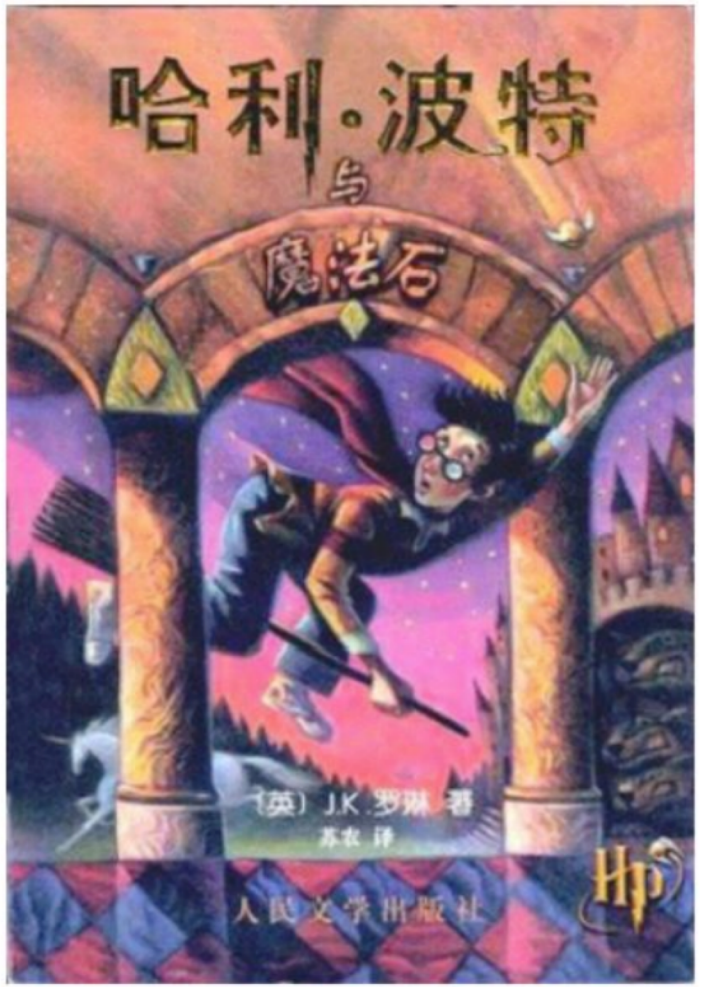

Parallel Corpora

				Sm.	Lg.
			CLASSIC SOUPS		
清燉雞湯	57.		House Chicken Soup (Chicken, Celery, Potato, Onion, Carrot)	1.50	2.75
雞飯湯	58.		Chicken Rice Soup	1.85	3.25
雞麵湯	59.		Chicken Noodle Soup	1.85	3.25
廣東雲吞	60.		Cantonese Wonton Soup	1.50	2.75
蕃茄蛋湯	61.		Tomato Clear Egg Drop Soup	1.65	2.95
雲吞湯	62.		Regular Wonton Soup	1.10	2.10
酸辣湯	63.	❧	Hot & Sour Soup	1.10	2.10
蛋花湯	64.		Egg Drop Soup	1.10	2.10
雲蛋湯	65.		Egg Drop Wonton Mix	1.10	2.10
豆腐菜湯	66.		Tofu Vegetable Soup	NA	3.50
雞玉米湯	67.		Chicken Corn Cream Soup	NA	3.50
蟹肉玉米湯	68.		Crab Meat Corn Cream Soup	NA	3.50
海鮮湯	69.		Seafood Soup	NA	3.50

Parallel Corpora (mining parallel data from microblogs Ling et al., 2013)

	ENGLISH	MANDARIN
1	i **wanna** live in a wes anderson world	我想要生活在Wes Anderson的世界里
2	Chicken soup, corn never truly digests. **TMI**.	鸡汤吧，玉米神马的从来没有真正消化过.恶心
3	To DanielVeuleman **yea iknw imma** work on that	对DanielVeuleman说，是的我知道，我正在向那方面努力
4	**msg 4** Warren G his **cday** is today 1 **yr** older.	发信息给Warren G，今天是他的生日，又老了一岁了。
5	Where **the hell** have you been all these years?	这些年你**TMD**到哪去了
	ENGLISH	ARABIC
6	It's **gonna** be a warm week!	الاسبو ع الياي حر
7	onni this gift only **4 u**	أوني هذة الهدية فقط لك
8	sunset in aqaba **:)**	غروب الشمس في العقبة:)
9	RT @MARYAMALKHAWAJA: there is a call for widespread protests in #bahrain **tmrw**	هناك نداء لظاهرات في عدة مناطق غدا

Table 2: Examples of English-Mandarin and English-Arabic sentence pairs. The English-Mandarin sentences were extracted from Sina Weibo and the English-Arabic sentences were extracted from Twitter. Some messages have been shorted to fit into the table. Some interesting aspects of these sentence pairs are marked in bold.

Discussions

- There is a lot more monolingual data in the world than translated data
- Easy to get about 1 trillion words of English by crawling the web
- With some work, you can get 1 billion translated words of English-French
  - What about Japanese-Turkish?

Phrase-Based MT

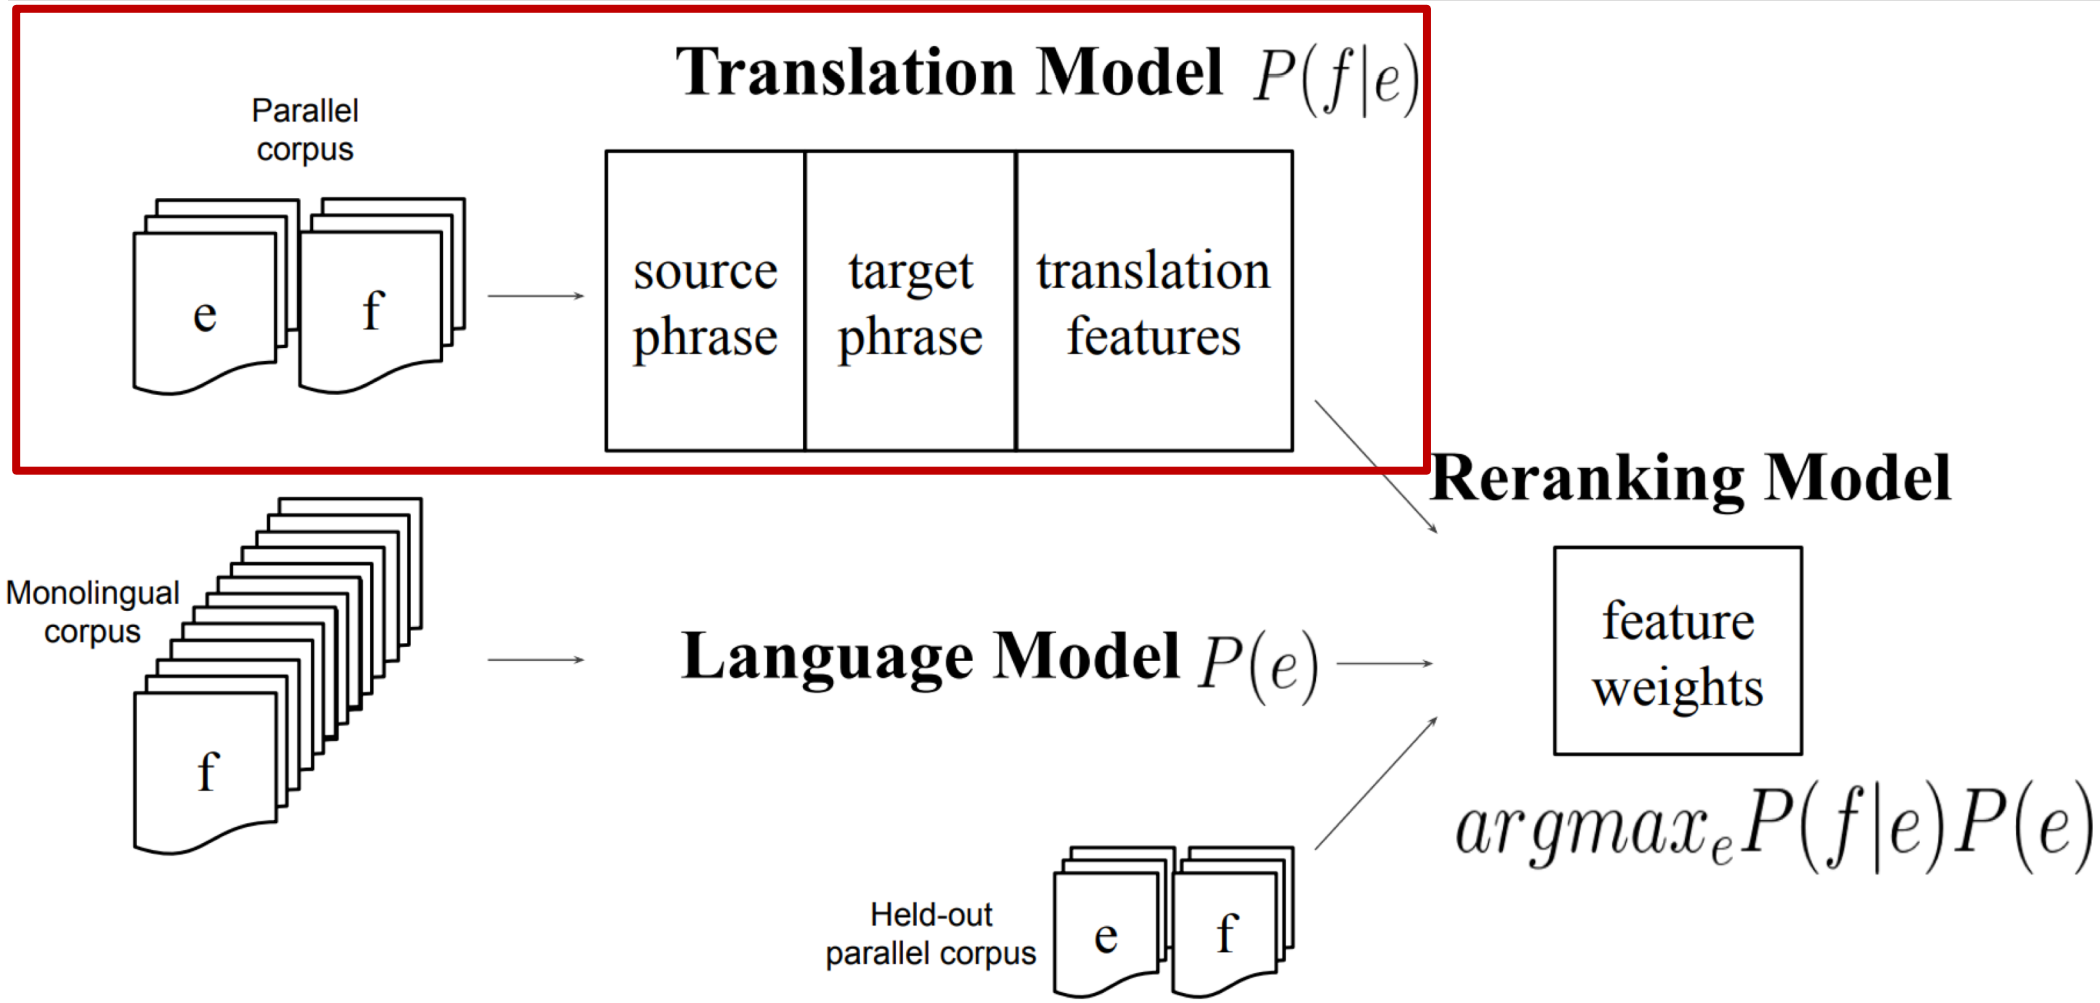

Construction of t-table

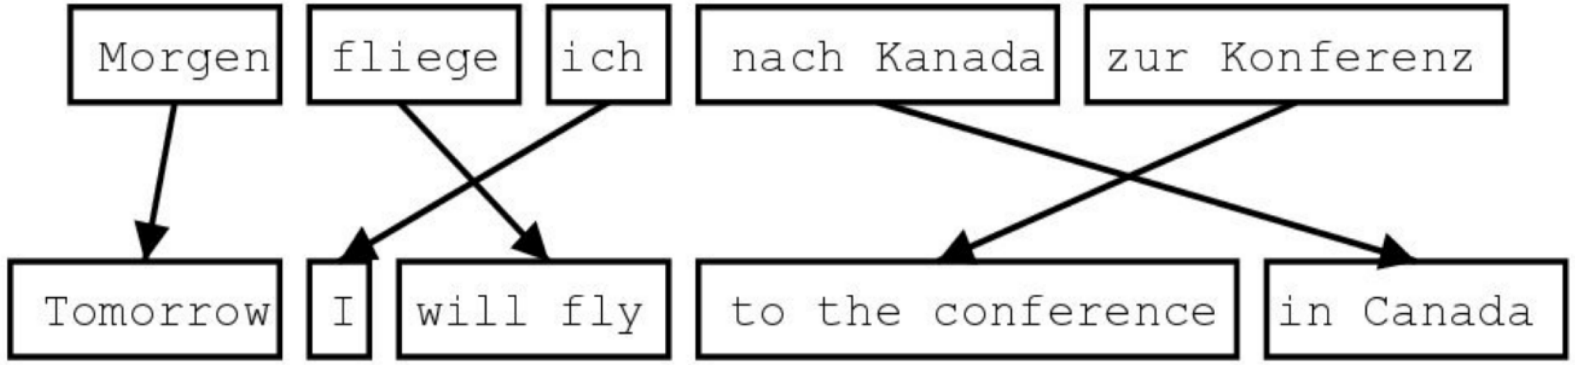

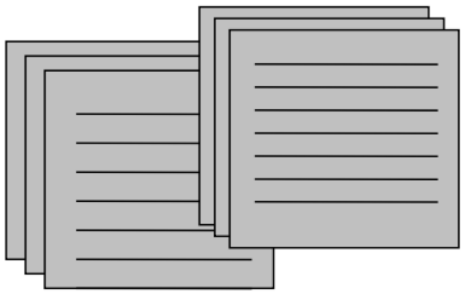

Sentence-aligned corpus

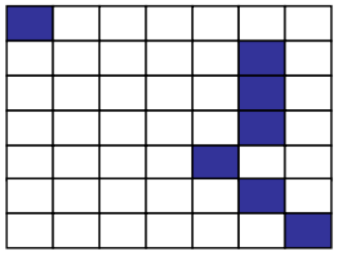

Word alignments

cat ||| chat ||| 0.9
the cat ||| le chat ||| 0.8
dog ||| chien ||| 0.8
house ||| maison ||| 0.6
my house ||| ma maison ||| 0.9
language ||| langue ||| 0.9
…

Phrase table (translation model)

Word Alignment Models

Lexical Translation

- How do we translate a word? Look it up in the dictionary

  Haus – house, building, home, household, shell

- Multiple translations
  - Some more frequent than others
  - Different word senses, different registers, different functions
  - *House, home* are common
- *Shell* is specialized (the Haus of a snail is a shell)

How Common is Each?

- Look at a parallel corpus (German text along with English translation)

Translation of Haus	Count
house	8000
building	1600
home	200
household	150
shell	50

Estimate Translation Probabilities

- Maximum likelihood estimation

$$\hat{p}_{\text{MLE}}(e \mid \texttt{Haus}) = \begin{cases} 0.8 & \text{if } e = \text{house}, \\ 0.16 & \text{if } e = \text{building}, \\ 0.02 & \text{if } e = \text{home}, \\ 0.015 & \text{if } e = \text{household}, \\ 0.005 & \text{if } e = \text{shell}. \end{cases}$$

Word Alignment

- Alignment can be visualized by drawing links between two sentences, and they are represented as vectors of positions

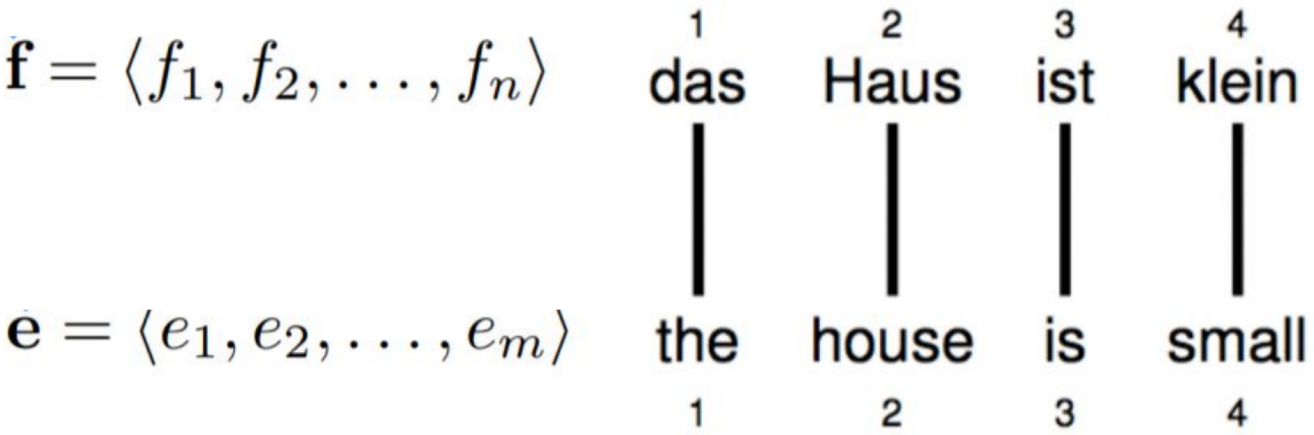

$$\mathbf{a} = (1, 2, 3, 4)^\top$$

Reordering

- Words may be reordered during translation

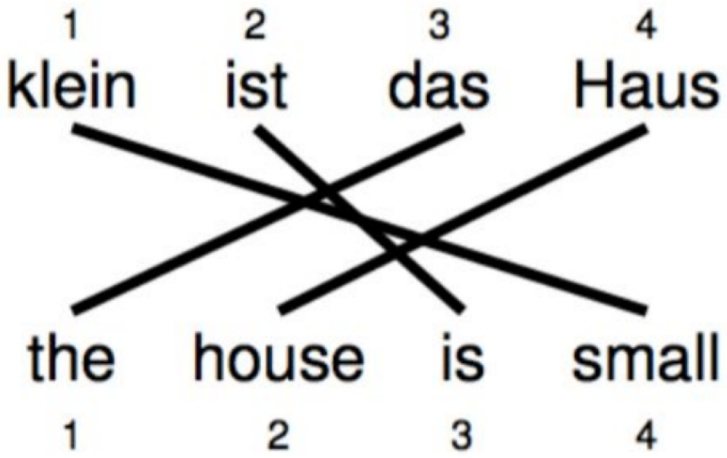

$$\mathbf{a} = (3, 4, 2, 1)^\top$$

Word Dropping

- A source word may not be translated at all

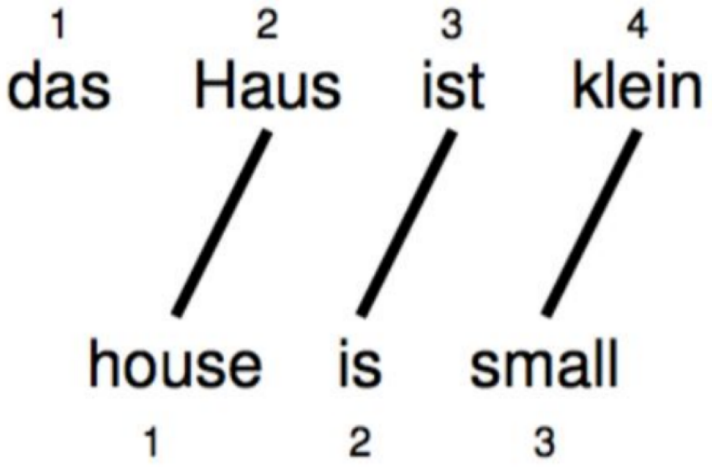

$$\mathbf{a} = (2, 3, 4)^\top$$

Word Insertion

- Words may be inserted during translation
  - English just does not have an equivalent
  - But it must be explained – we typically assume every source sentence contains a NULL token

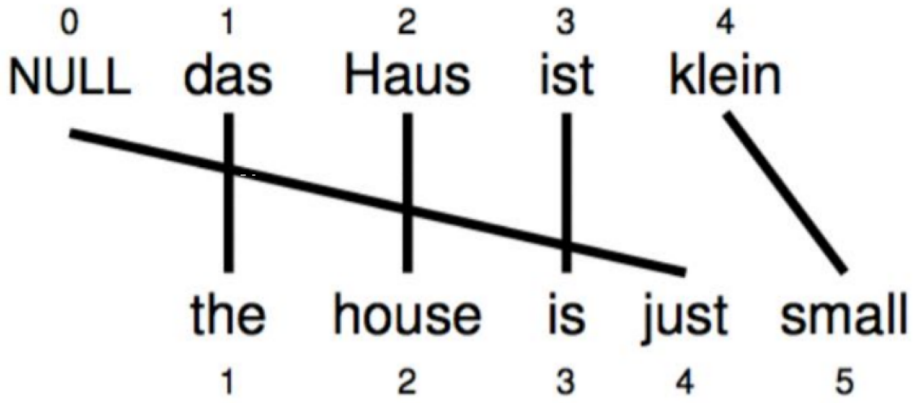

$$\mathbf{a} = (1, 2, 3, 0, 4)^\top$$

One-to-many Translation

- A source word may translate into more than one target word

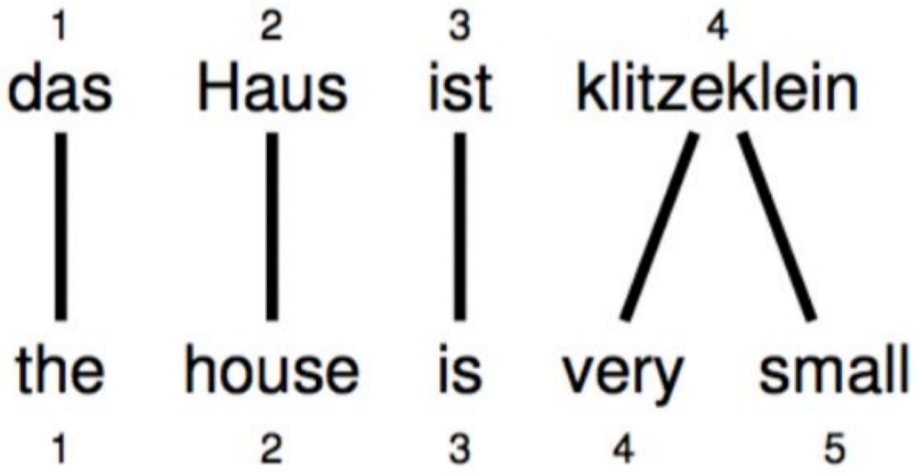

$$\mathbf{a} = (1, 2, 3, 4, 4)^\top$$

Many-to-one Translation

- More than one source word may not translate as a unit in lexical translation

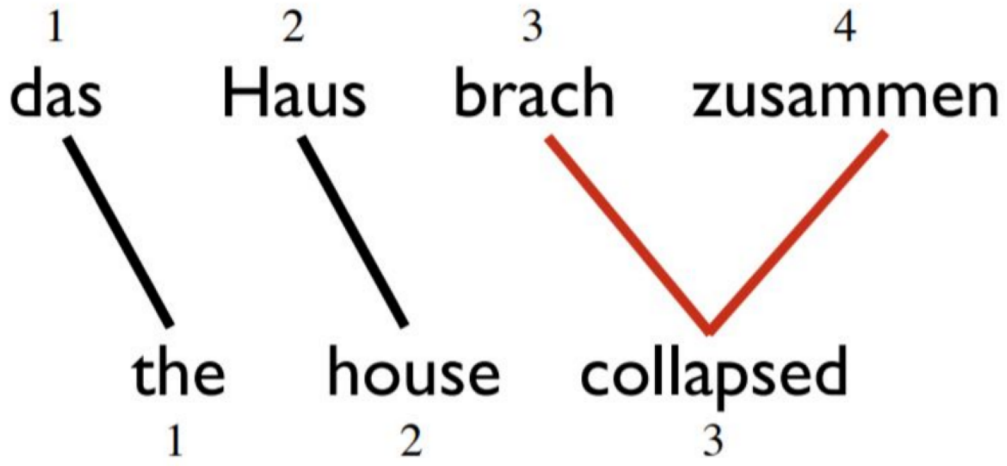

$\mathbf{a} = ???$ $\mathbf{a} = (1, 2, (3, 4)^\top)^\top$?

Computing Word Alignments

- Word alignments are the basis for most translation algorithms
- Given two sentences F and E, find a good alignment
- But a word-alignment algorithm can also be part of a mini-translation model itself
- One the most basic alignment models is also a simplistic translation model